

Authentication - you will need to upload the service account key json file when prompted

```
from google.colab import auth
auth.authenticate_service_account()
```

➞ Successfully saved credentials for pipelineauth@weatherlink-404323.iam.gserviceaccount.com

+ Code

+ Text

Testing querying data from GBQ now that we are authenticated:

```
from google.cloud import bigquery
project_name='weatherlink-404323'
client = bigquery.Client(project=project_name)
```

```
dataset_name = 'weatherlink_master'
table_name = 'census'
```

```
sql_query = (f"SELECT * FROM {dataset_name}.{table_name} LIMIT 10")
```

```
df = client.query(sql_query).to_dataframe()
```

```
print(df)
```

	geo_id	income_per_capita	median_income	year
0	36079	106871	48417	2021
1	47065	66096	38668	2021
2	12097	60585	26789	2021
3	9003	80069	43636	2021
4	42017	100144	50607	2021
5	34029	75719	39055	2021
6	18089	61443	31785	2021
7	6067	80063	37259	2021
8	49011	93182	38879	2021
9	22071	46942	35587	2021

Analyzing some of the census data, I want to see what the number of columns in each year are to see if they differ.

```

accident_client = bigquery.Client()

for i in range(2007,2024):
    table_name = f"bigquery-public-data.census_bureau_acs.county_{i}_1yr"

    sql_query = (f"SELECT * FROM {table_name} LIMIT 1")

    try:
        df = accident_client.query(sql_query).to_dataframe()
        print (f"Table name {table_name} has {(df.shape[1])} cols")

    #print(df)
    except:
        print (f"Table name {table_name} was not found :( ")

```

```

Table name bigquery-public-data.census_bureau_acs.county_2007_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2008_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2009_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2010_1yr has 247 cols
Table name bigquery-public-data.census_bureau_acs.county_2011_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2012_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2013_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2014_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2015_1yr has 246 cols
Table name bigquery-public-data.census_bureau_acs.county_2016_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2017_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2018_1yr has 252 cols
Table name bigquery-public-data.census_bureau_acs.county_2019_1yr has 245 cols
Table name bigquery-public-data.census_bureau_acs.county_2020_1yr was not four

```

Okay, now I want to see WHICH columns everything does not have, so I can see what we would be missing out on if we just ignore columns that are not common across all data sets.

```

import pandas as pd

year_start = 2007
year_stop = 2025

# Initialize a dictionary to store the schema for each table
table_schemas = {}

```

```

# Initialize a set to store common columns
common_columns_set = None

# Initialize a dictionary to store unique columns for each table
unique_columns_dict = {}

# Initialize a list to store DataFrames for each year
dfs = []

for i in range(year_start, year_stop):
    table_name = f"bigquery-public-data.census_bureau_acs.county_{i}_1yr"

    try:
        # Fetch the schema (column information) for each table
        if table_name not in table_schemas:
            table = accident_client.get_table(table_name)
            table_schemas[table_name] = set([field.name for field in table.schema])

        # If this is the first DataFrame, initialize the set with its columns
        if common_columns_set is None:
            common_columns_set = set(table_schemas[table_name])
        else:
            # Update the set to include only columns present in both DataFrames
            common_columns_set.intersection_update(table_schemas[table_name])

        # Update the set to include only columns not present in other DataFrames
        unique_columns_dict[table_name] = table_schemas[table_name].difference(common_columns_set)

        print(f"Table name {table_name} has {len(table_schemas[table_name])} cols")

        # Fetch the data and add it to a DataFrame
        query = f"SELECT * FROM {table_name}"
        df = accident_client.query(query).to_dataframe()

        # Add the 'year' column to the DataFrame
        df['year'] = i

        # Keep only the common columns
        df = df[list(common_columns_set) + ['year']]

        # Append the DataFrame to the list
        dfs.append(df)

    except:
        print(f"Table {table_name} not found")

# Print unique columns for each table
for i in range(year_start, year_stop):

```

```

table_name = f"bigquery-public-data.census_bureau_acs.county_{i}_1yr"
try:
    print(f"Unique columns in {table_name}: {unique_columns_dict[table_name]}")
except:
    print(f"Table {table_name} not found")

# Concatenate all DataFrames into a master DataFrame
master_df = pd.concat(dfs, ignore_index=True)

# Print the master DataFrame
print("\nMaster DataFrame:")
print(master_df)

```

```

Table name bigquery-public-data.census_bureau_acs.county_2007_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2008_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2009_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2010_1yr has 247 cc
Table name bigquery-public-data.census_bureau_acs.county_2011_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2012_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2013_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2014_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2015_1yr has 246 cc
Table name bigquery-public-data.census_bureau_acs.county_2016_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2017_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2018_1yr has 252 cc
Table name bigquery-public-data.census_bureau_acs.county_2019_1yr has 245 cc
Table bigquery-public-data.census_bureau_acs.county_2020_1yr not found
Table name bigquery-public-data.census_bureau_acs.county_2021_1yr has 245 cc
Table bigquery-public-data.census_bureau_acs.county_2022_1yr not found
Table bigquery-public-data.census_bureau_acs.county_2023_1yr not found
Table bigquery-public-data.census_bureau_acs.county_2024_1yr not found
Unique columns in bigquery-public-data.census_bureau_acs.county_2007_1yr: se
Unique columns in bigquery-public-data.census_bureau_acs.county_2008_1yr: se
Unique columns in bigquery-public-data.census_bureau_acs.county_2009_1yr: se
Unique columns in bigquery-public-data.census_bureau_acs.county_2010_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2011_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2012_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2013_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2014_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2015_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2016_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2017_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2018_1yr: {'
Unique columns in bigquery-public-data.census_bureau_acs.county_2019_1yr: {'
Table bigquery-public-data.census_bureau_acs.county_2020_1yr not found
Unique columns in bigquery-public-data.census_bureau_acs.county_2021_1yr: {'
Table bigquery-public-data.census_bureau_acs.county_2022_1yr not found
Table bigquery-public-data.census_bureau_acs.county_2023_1yr not found
Table bigquery-public-data.census_bureau_acs.county_2024_1yr not found

```

```

Master DataFrame:
    income_per_capita  two_cars  pop_5_years_over  occupation_services  \

```

0	24635.0	94225.0	NaN	48502.0
1	22853.0	68027.0	NaN	37994.0
2	37651.0	125047.0	NaN	74770.0
3	21887.0	50863.0	NaN	60063.0
4	27268.0	106823.0	NaN	62354.0
...
11503	41636.0	55639.0	NaN	39743.0
11504	30473.0	60722.0	NaN	35711.0
11505	20943.0	51464.0	NaN	38167.0
11506	37681.0	89537.0	NaN	42335.0
11507	35375.0	113759.0	NaN	55175.0

	owner_occupied_housing_units	female_75_to_79	no_car	pop_divorced	\
0	162734.0	7125.0	6003.0	48781.0	
1	130322.0	7678.0	5750.0	40576.0	
2	213786.0	7548.0	31819.0	83851.0	
3	119820.0	8177.0	44853.0	50839.0	
4	189723.0	10108.0	6811.0	75202.0	
...

Okay, so we can see there are indeed some columns that are NOT in all of the data sets. We now have a master dataframe of ALL the census data, with only the columns from every set. (There is a gap from the year 2020, where no census data was generated due to the COVID-19 pandemic)

Since 253! has 500 digits, we are not going to be making scatter plots of every column vs every other column. Instead, we are going to upload this dataframe to Google Big Query in our 'Data Warehouse', and then select a number of columns to compare for analysis.

```
pip install pandas_gbq
```

```

from pandas_gbq import to_gbq
from pandas_gbq.schema import generate_bq_schema
from io import StringIO

client = bigquery.Client()

project_id='weatherlink-404323'
dataset_id = 'weatherlink_master'
table_id = 'census_master'

# This is a weird work around to get the dataframe acceptable for upload
# temporarily store the dataframe as a csv in a string variable
temp_csv_string = master_df.to_csv(sep=";", index=False)
temp_csv_string_I0 = StringIO(temp_csv_string)
# create new dataframe from string variable
new_df = pd.read_csv(temp_csv_string_I0, sep=";")

to_gbq(new_df, f"{dataset_id}.{table_id}", project_id=project_id, if_exists='repl

```

```

100%|██████████| 1/1 [00:00<00:00, 1348.22it/s]

```

Lets test by querying our new master table

```

sql = f" SELECT COUNT(*) FROM {project_id}.{dataset_id}.{table_id}"

result = client.query(sql)

print(result)

```

```

QueryJob<project=weatherlink-404323, location=US, id=a28bd2f3-62ce-40c4-9140-!

```

